The design leaves

One document per domain of the architecture, plus the schema, the API and the migration. The trunk's document map shows where each leaf hangs and which legacy documents it builds on.

Trunk section	Leaf	Covers
6.1	SOURCES.md	every external source behind an adapter; bar feeds and the cache; fallback orders; health; keys
6.2	TOOLS-AND-SCREENERS.md	every tool and screener rule by rule; the builder; discovery; lifecycle
6.3	CARD.md	the card: all 97 legacy fields, fact families, the builder, presentation
6.4	MARKET.md	the market read: regime, sectors, hot sectors, M and L on the tab
6.5	ONEIL.md	the O'Neil chain: every letter's rules, what was built, the research worker
6.6	REGISTERS.md	R0–R4 and the shortlist, every column; the research record
6.7	LEARNING.md	the score and the scorecard
6.8	QUANT.md	the library, strategies, backtests, alerts, print sheets
6.9	JOURNAL.md	every real trade: accounts, fills, pairing, stops, the page
6.10	PLATFORM.md	login, scheduler, settings, health, backup, restore, deploy, the page shell
5	DATA-MODEL.md	every table, and where every legacy table goes
4	API.md	every route
10	MIGRATION.md	the 14 steps, their checks and undo, the import of the old data

Each domain leaf has the same five parts:

1. **Its place in the trunk**: the section it details, and the domains it talks to.
2. **The legacy design, in full**: every field, source, calculation and rule the old desk had, so nothing it knew is lost.
3. **What failed**: the faults found in the 2026-10 review, by ID.
4. **The new design**: how it is built, and why.

5. **Contracts**: the tables, fact families and routes it owns or reads, and how the rebuild is proven.