The desk: one running tool, end to end

Above: DESK.md §4 (running tools) and §4.5 (the research each tool runs).

Talks to:

- the web layer, which brings its pages and API;
- the scheduler, which runs its scans, refreshes, captures and outcomes in its context;
- the market reader, for one market read per scan;
- the data sources (candidates, quotes, short interest, borrow, bars) and the research worker (the O'Neil letters);
- the shared record and its own file, of which it is the only writer.

---

1. What it is

One tool, as one experiment. A tool is one way of finding stocks: its own screeners, its own cards, its own daily record, its own shortlist, its own score, kept apart from every other tool (DESK.md §3). It finds the stocks its screeners match, turns each into a card with everything known about it, keeps today's live list, freezes a photo of that list once a day, measures how far each photographed stock moved, scores its cards, and serves its pages.

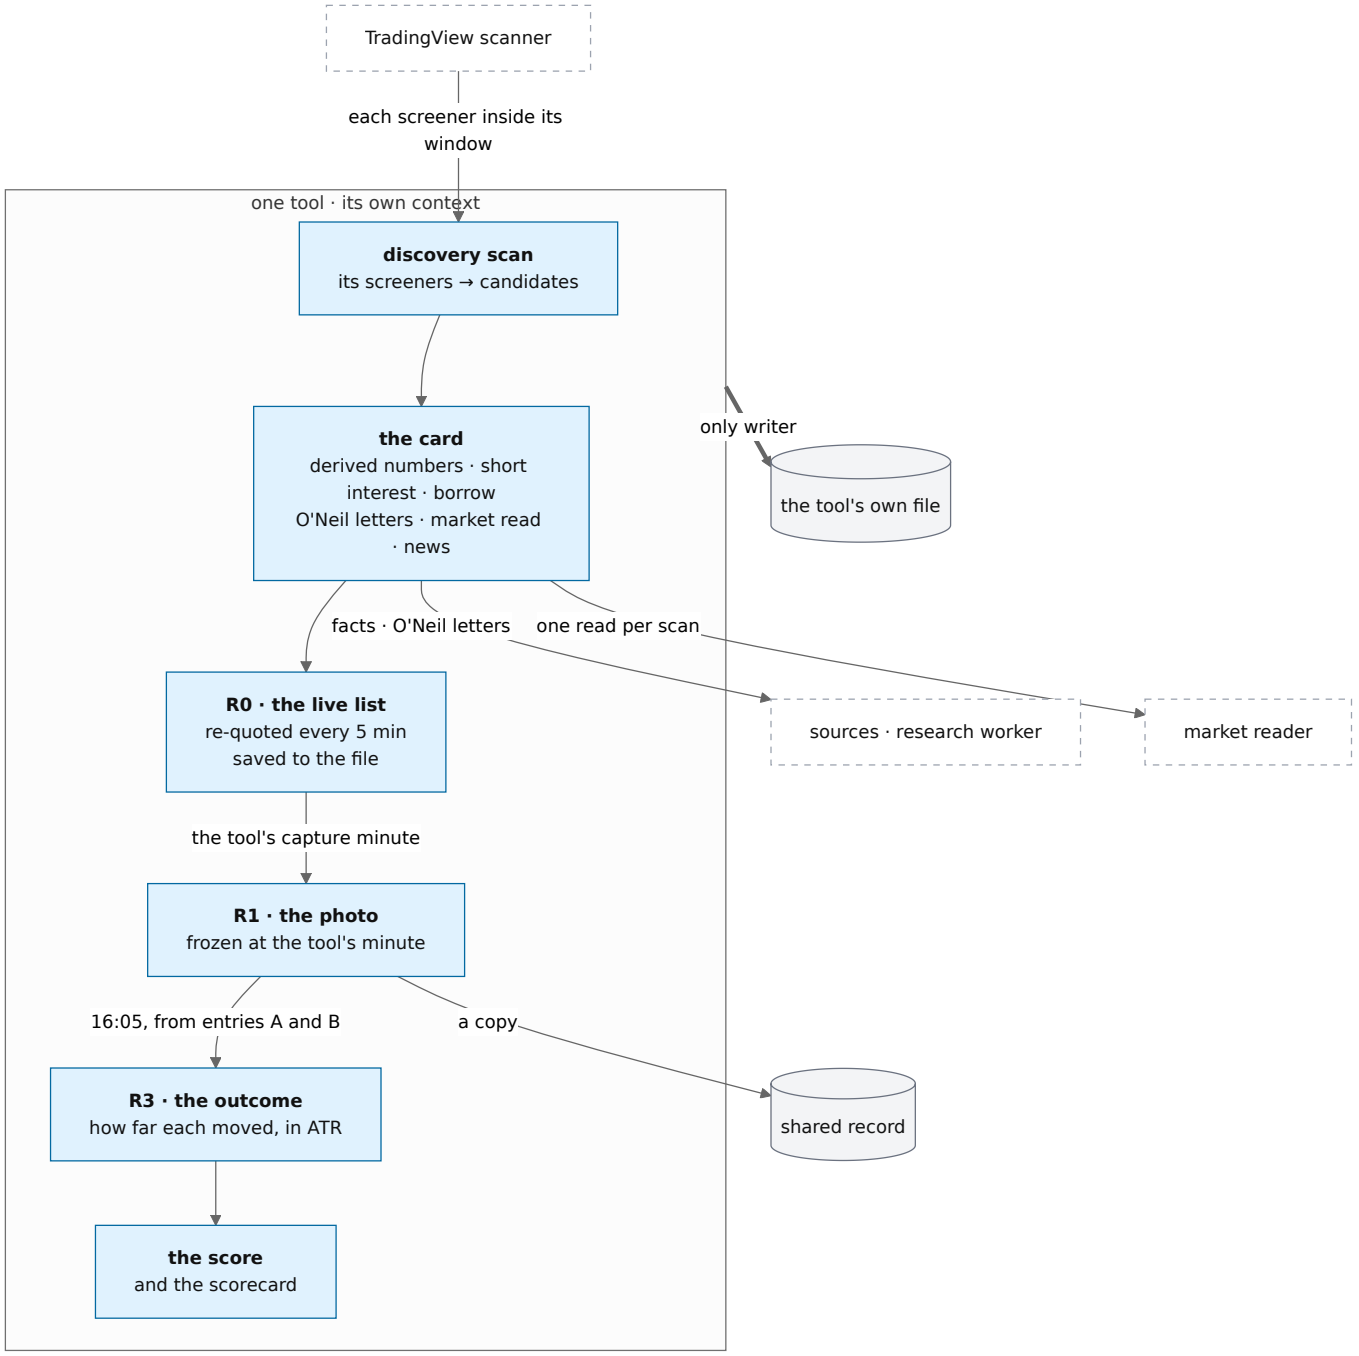
This document is the shape of a running tool end to end; each part has its own document one level down:

Part	Document
its screeners and how it finds stocks	tool/SCREENERS.md
the card: every field and its source	tool/CARD.md
its daily record and the outcomes	tool/REGISTERS.md
the score and the scorecard	tool/LEARNING.md
its pages	tool/PAGES.md

What it owns: its screeners, its live list, its registers, its shortlist and its score — all in its own file. **What it never does:** see or change another tool's results, compute O'Neil data (the worker does), or place an order.

---

2. The diagram



3. Its parts, in order

3.1 Find: the screeners and the scan

Each tool asks **TradingView's scanner** one request per enabled screener whose run window contains the current minute. Every request carries the screener's own filters **plus the tradability floor** (a price, average-volume and ATR minimum applied to every screener, with named exceptions) and a base stock filter. A changed or misaligned answer is refused rather than shifting every field by one column (DATA-SOURCES.md). The matches of every screener are

merged into one row per ticker, which screener keys list which screeners found it. Detail: tool/SCREENERS.md.

3.2 Know: the card

Each row becomes a card: derived numbers (gap, ranges, ADR %, distances to each level), short interest and borrow, the O'Neil letters (read from the research worker, not computed here), the market read stamped on from the market reader, and news with a catalyst. A missing value says it is missing and why — never 0, never bearish. Every field, its source, and when a blank is legitimate: tool/CARD.md.

3.3 Keep: the live list (R0)

The cards of today, live (matched this scan) and stale (matched earlier, kept on screen and re-quoted), held in the tool's context and **saved to its file after every change**, so a restart brings the list back rather than emptying the screen for the rest of the day (DESK.md §7). Cleared at midnight.

3.4 Record: the photo (R1) and the market (R2)

At the tool's **own minute** in the first hour (09:36 … 10:16 today), every card on the list is frozen exactly as known then — the photo. **Only the photo judges the tool.** A missed minute (the desk was down) is marked, never taken late as if on time. The market reader stores R2, the market at fixed minutes around the photo. Detail: tool/REGISTERS.md.

3.5 Measure: the outcome (R3)

After the close, from two fixed entry minutes (A and B, a few minutes after the photo), the desk measures how far each photographed stock moved up and down to the close, in units of the stock's average daily range (ATR). **A move of 1.3 ATR or more is a good move.** There is no stop, target or exit: this judges what a tool finds, not a trade.

3.6 Learn: the score and the scorecard

Over weeks, the photos and their outcomes show **which screeners find stocks that move** — judged by the move, with consistency across days deciding, against the tool's own baseline, the **benchmark** tool (which every other must beat by a clear margin) and the **mirror** of each screener (the same setup facing the other way, to test whether its direction is real). The order of work: collect for about a month, cut the screeners whose stocks do not move, then work on entries. A model to rank a tool's stocks is prepared, not built, until the trader chooses one. Detail: tool/LEARNING.md.

3.7 Show: its pages

The tool's eight tabs (its cards, the market read, the shortlist, the warehouse of registers, the scorecard, the screener builder, settings, the monitor), served in its context on its own port (WEB.md). Detail: tool/PAGES.md.

---

4. Contracts

It owns (the tool's own file, SQLite — the only writer)	Holds
its screeners and their versions	the definitions, each change a new version
R0 (the live list), R1 (the photos), R3 (the outcomes), the training rows, the shortlist	one file per tool (FILES.md, tool/REGISTERS.md)
its settings	the tradability floor, capture times, the shortlist rule, the scorer switch

It reads	From
candidates, quotes, levels, ranges, short interest, borrow, bars	the data sources (DATA-SOURCES.md)
the O'Neil letters for its tickers	the research worker (RESEARCH-WORKER.md)
the market read	the market reader (MARKET.md), once per scan

It copies	To
each photo (R1) and its outcomes	the shared record, so qp and the journal read its history without opening its file (SHARED-RECORD.md)

Its jobs — discovery, refresh, the shortlist rule, the photo, the outcomes, the scorer training — run in its context on the one scheduler (JOBS.md).

5. When it fails

Event	What happens
a screener's answer is changed or misaligned	that answer is refused and the scan report says why; no card gets shifted values
a source fails	that field shows `unavailable` with the reason; the card still builds
the desk is down at the tool's photo minute	the minute is marked missed and the phone is told; never taken late
a scan outlasts its slot	the next slot is skipped and recorded; two scans of one tool never run at once
the tool is paused	scans return at once saying so; the cards already found stay; the registers are untouched

6. Built on

The old desk: every tool was its own program, running this pipeline against its own screeners. The pipeline is legacy/screeners/DATA-PIPELINE.md; the method,

legacy/screeners/METHODOLOGY.md; the registers, legacy/screeners/REGISTERS.md; the lifecycle, legacy/screeners/TOOL-LIFECYCLE.md.

What failed there (the tool-level faults; the desk-level ones are in DESK.md §7):

ID	Fault	Prevented by
—	the live list lived only in memory and was lost on restart	saved to the tool's file after every change
D-S12	the scorecard graded every screener on the upward move, so a bearish screener's verdict was backwards	judged by the move in the screener's own direction (tool/LEARNING.md)
D-D1	with the scorer off, the shortlist auto-rule added nothing	the score and the shortlist rule (tool/LEARNING.md)
F-14	a missing input shown as zero	a missing input gives a missing output (tool/CARD.md)

7. Status

Part	Status
the pipeline (screeners, card, live list, photo, outcomes), as each tool's own program	**built** (legacy)
the same pipeline run in the desk's context, tool by tool	**to build** (step B), each only after a clean week of photos (DESK.md §3, the comparison copy)
the score and scorecard rebuilt	**to build** (step E–F; tool/LEARNING.md)
the level-3 documents (SCREENERS, CARD, REGISTERS, LEARNING, PAGES)	**to write**